Two quantization experiments

INT8 and INT4 · two ways to use reduced precision

Weights only

Weights saved as INT8 or INT4.

Calculations run in FP32.

Weights + calculations

Weights and inputs are integers.

Matrix products run with integers.

Softmax, sigmoid and some other steps remain in FP32.