Integer quantization · stored model tensors (KiB)

Dataset	Model	FP32	INT8	INT4
OPS-SAT	TranAD+	1.1	0.8	0.7
	One-decoder	0.7	0.5	0.5
	One-pass	0.7	0.5	0.5
	Encoder-only	0.3	0.3	0.3
SMAP	TranAD+	265.9	77.0	44.7
	One-decoder	166.3	48.7	28.5
	One-pass	166.3	48.7	28.5
	Encoder-only	66.7	20.4	12.4
MSL	TranAD+	1285.4	357.4	200.8
	One-decoder	808.5	230.6	133.1
	One-pass	808.5	230.6	133.1
	Encoder-only	331.5	103.9	65.5

Reference W/FF · packed tensors, biases, scales and buffers included · training checkpoints excluded